

ADVERTISING USING ELASTICNET

In []:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn import preprocessing, svm
```

In []:

```
df=pd.read_csv(r"/content/Advertising.csv")
df
```

Out[]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In []:

```
df.head()
```

Out[]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In []:

```
df.tail()
```

Out[]:

	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6

	TV	Radio	Newspaper	Sales
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

In []:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   TV          200 non-null    float64
 1   Radio        200 non-null    float64
 2   Newspaper    200 non-null    float64
 3   Sales         200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

In []:

```
df.describe()
```

Out[]:

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

In []:

```
df.shape
```

Out[]:

```
(200, 4)
```

In []:

```
df.columns
```

Out[]:

```
Index(['TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')
```

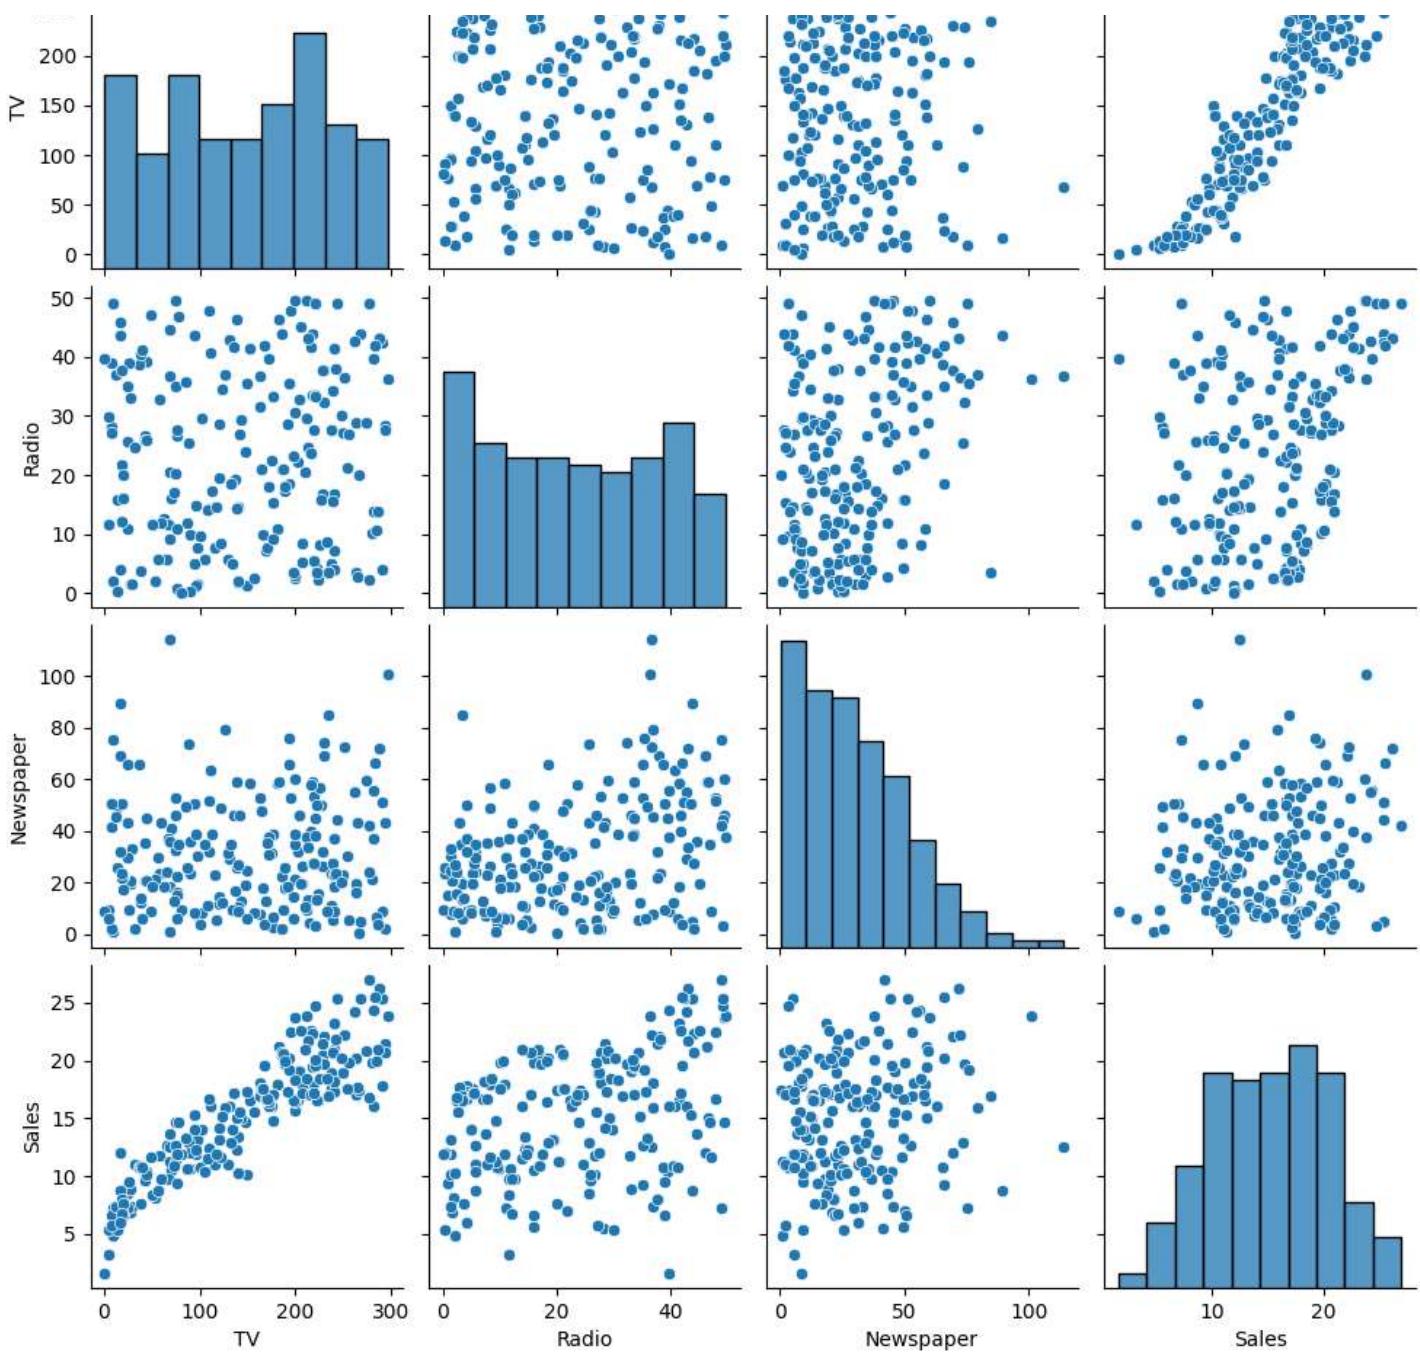
In []:

```
#EDA
sns.pairplot(df)
```

Out[]:

```
<seaborn.axisgrid.PairGrid at 0x7f0860b76bf0>
```



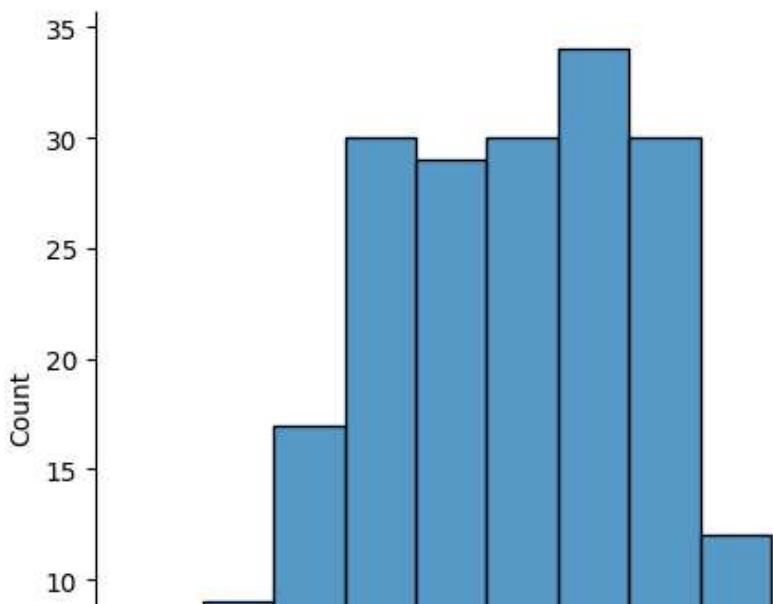


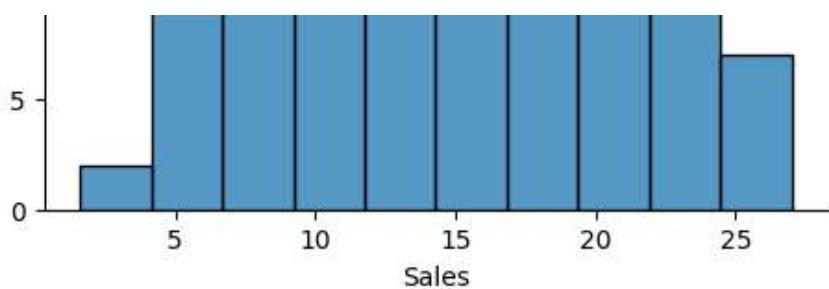
In []:

```
sns.displot(df['Sales'])
```

Out[]:

```
<seaborn.axisgrid.FacetGrid at 0x7f085fe048b0>
```



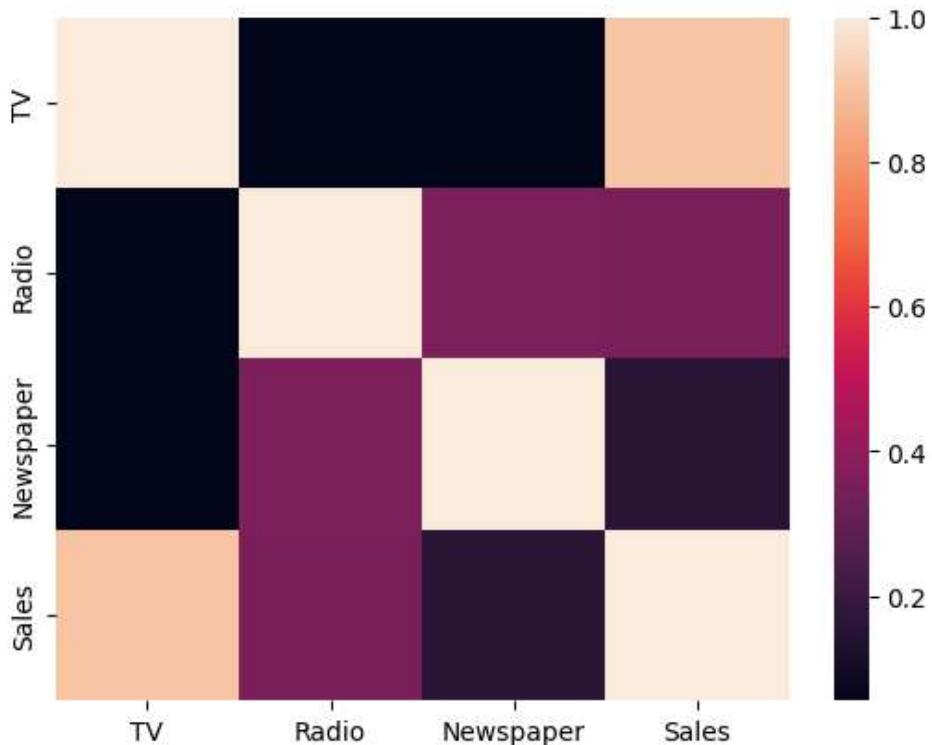


In []:

```
addf=df[['TV', 'Radio', 'Newspaper', 'Sales']]
sns.heatmap(addf.corr())
```

Out[]:

<Axes: >



In []:

```
X=addf[['TV', 'Radio', 'Newspaper']]
y=df['Sales']
```

In []:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=101)
from sklearn.linear_model import LinearRegression
lm=LinearRegression()
lm.fit(X_train,y_train)
print(lm.intercept_)
```

4.681232151484295

In []:

```
coeff_df=pd.DataFrame(lm.coef_,X.columns,columns=['coefficient'])
coeff_df
```

Out[]:

coefficient

TV	0.054930
----	----------

Radio 0.109558
Newspaper -0.006194

In []:

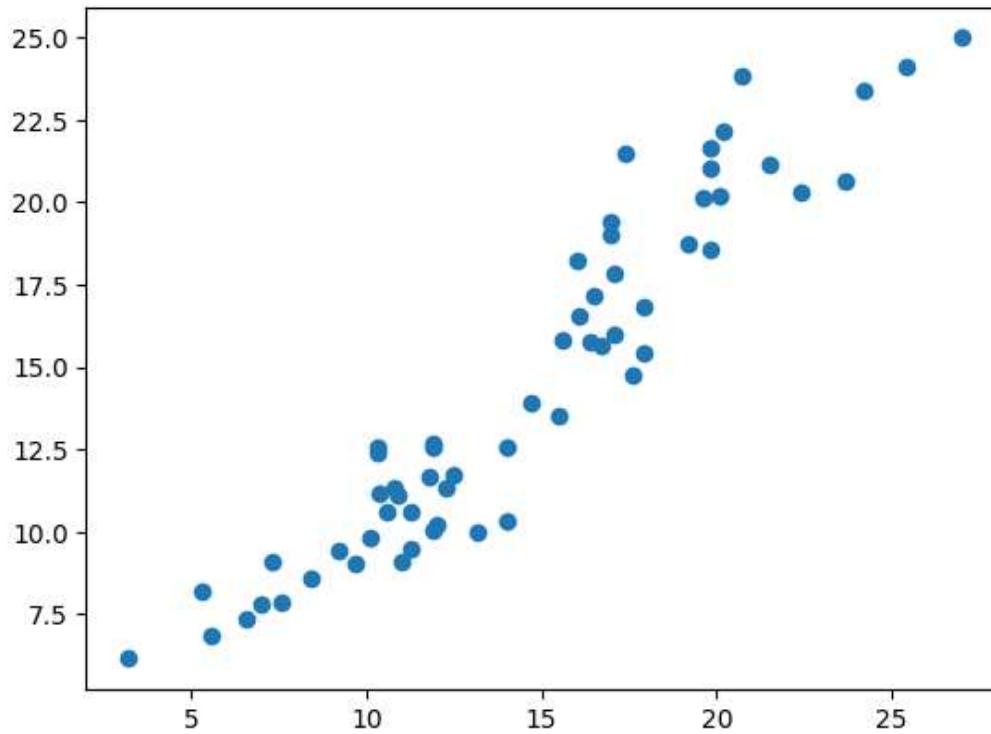
```
predictions=lm.predict(X_test)
```

In []:

```
plt.scatter(y_test,predictions)
```

Out[]:

```
<matplotlib.collections.PathCollection at 0x7f085d916650>
```

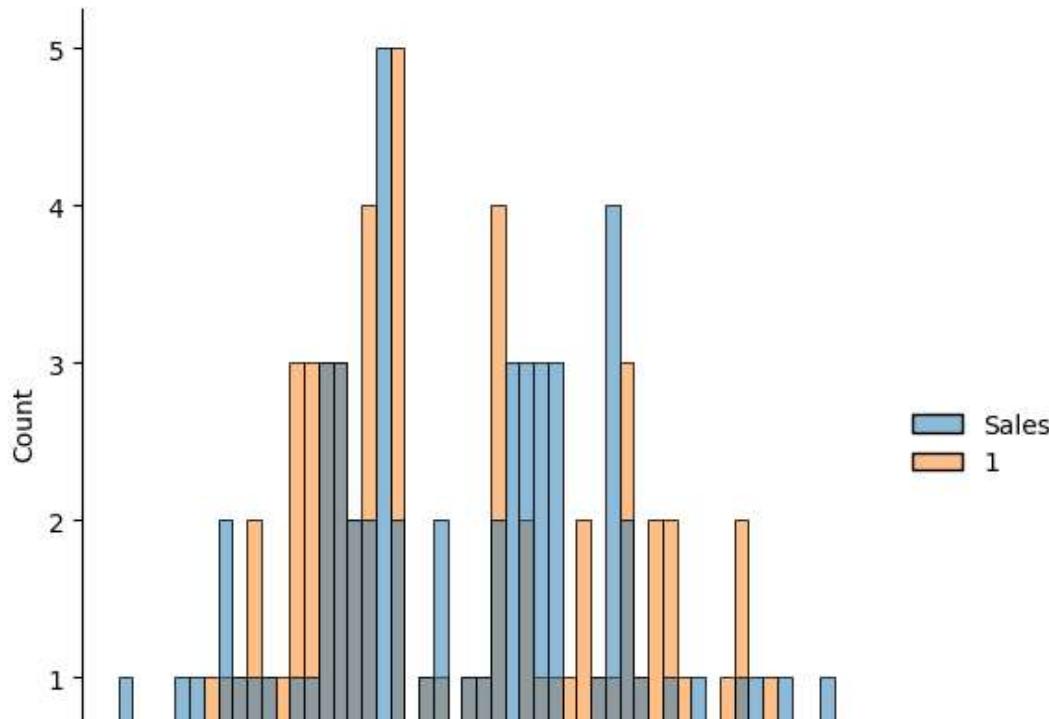


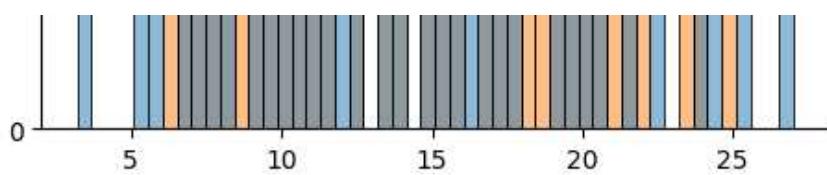
In []:

```
sns.displot((y_test,predictions),bins=50) #without semicolon
```

Out[]:

```
<seaborn.axisgrid.FacetGrid at 0x7f085d8e6ce0>
```





In []:

```
from sklearn import metrics
print('MAE:',metrics.mean_absolute_error(y_test,predictions))
print('MSE:',metrics.mean_squared_error(y_test,predictions))
print('RMSE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

MAE: 1.3731200698367851

MSE: 2.8685706338964962

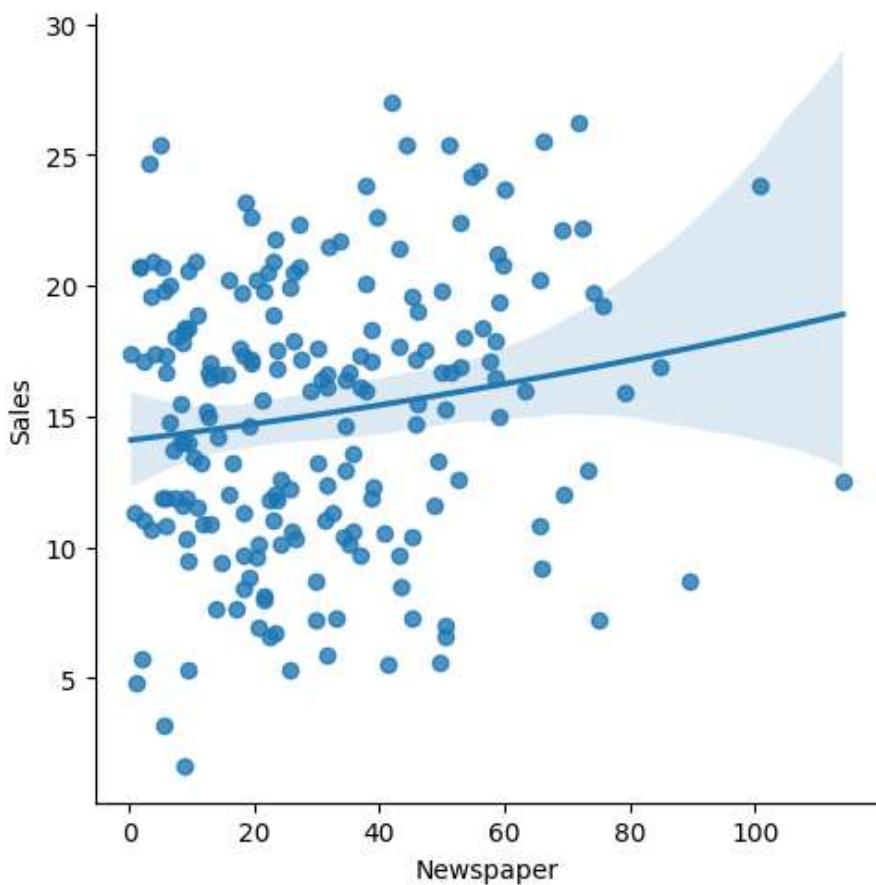
RMSE: 1.6936855180040054

In []:

```
sns.lmplot(x="Newspaper",y="Sales",data=df,order=2)
```

Out[]:

```
<seaborn.axisgrid.FacetGrid at 0x7f08592bd480>
```



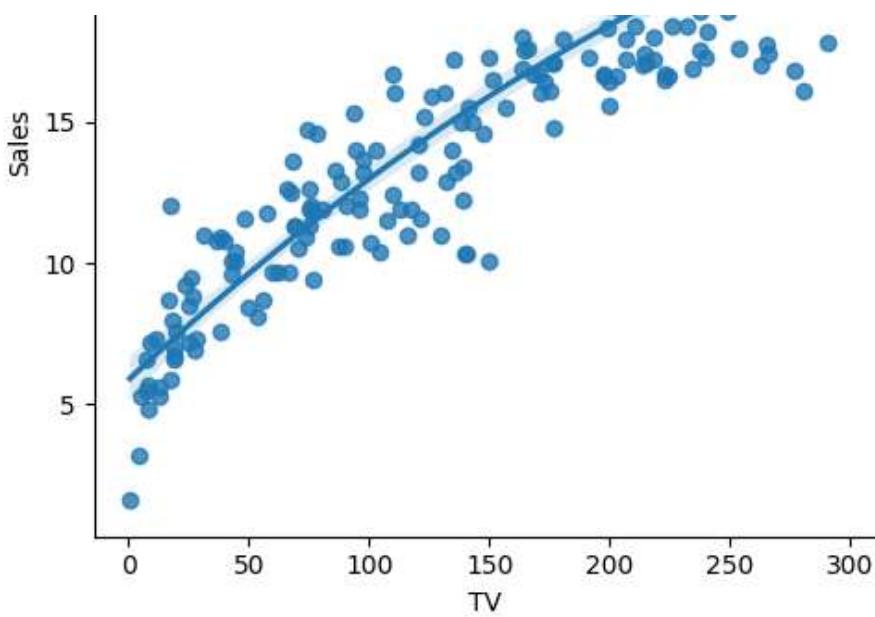
In []:

```
sns.lmplot(x="TV",y="Sales",data=df,order=2)
```

Out[]:

```
<seaborn.axisgrid.FacetGrid at 0x7f085d965d20>
```



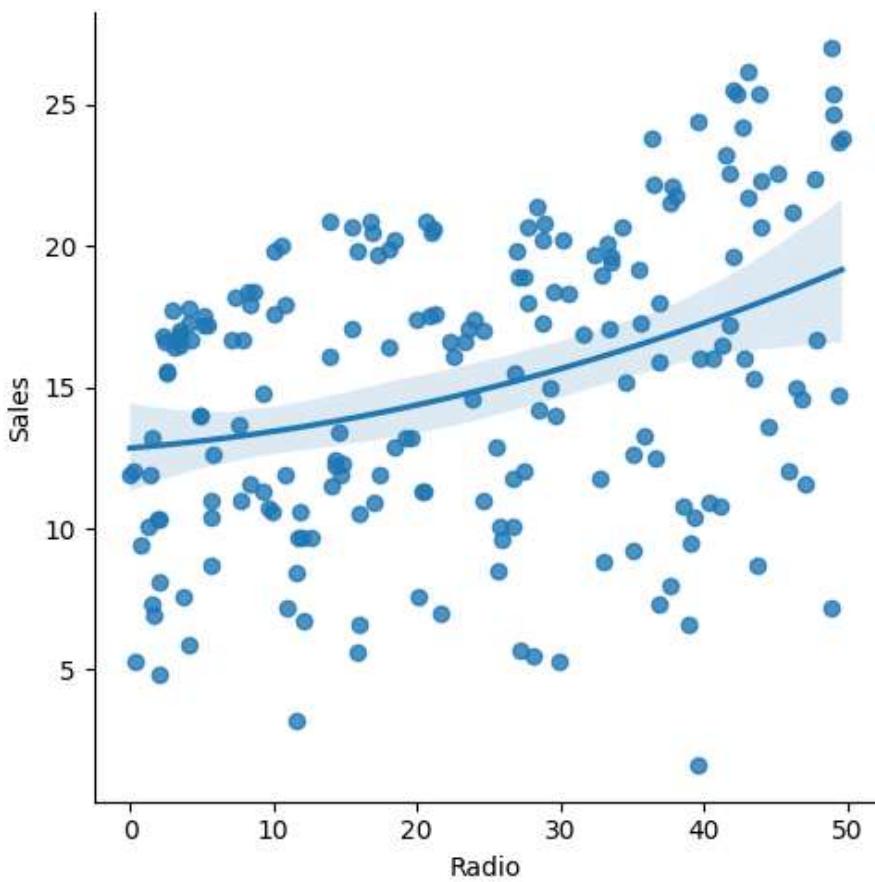


In []:

```
sns.lmplot(x="Radio", y="Sales", data=df, order=2)
```

Out[]:

```
<seaborn.axisgrid.FacetGrid at 0x7f08593d3b50>
```



In []:

```
df.fillna(method='ffill', inplace=True)
```

In []:

```
regr=LinearRegression()
```

In []:

```
x=np.array(df['TV']).reshape(-1,1)
y=np.array(df['Sales']).reshape(-1,1)
```

```
df.dropna(inplace=True)
```

In []:

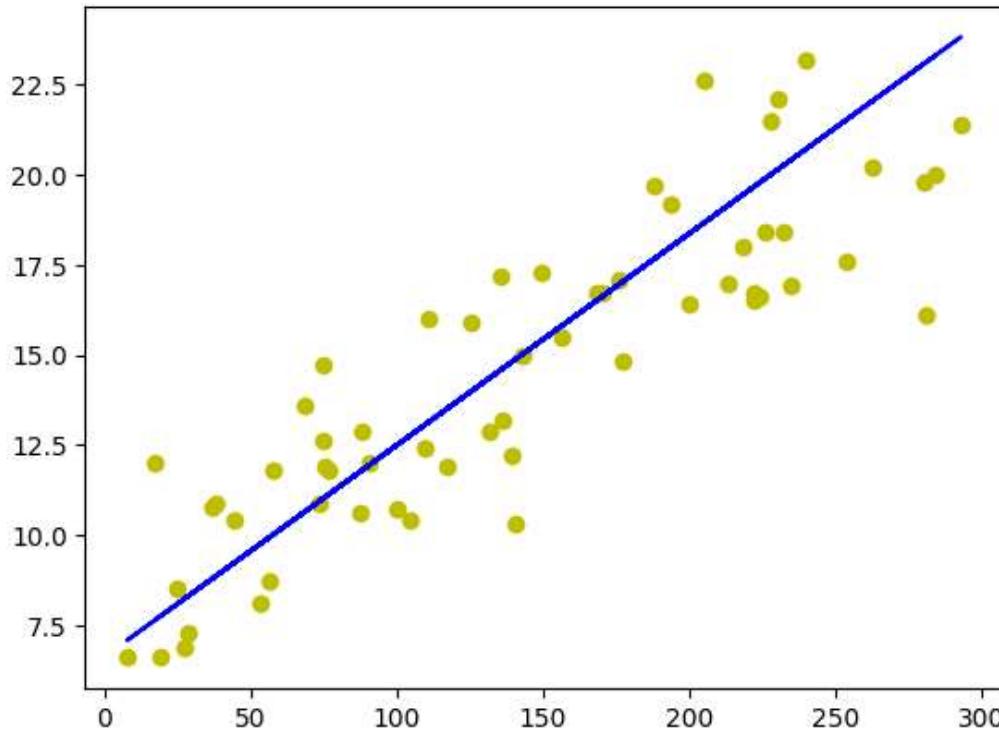
```
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
```

Out[]:

```
▼ LinearRegression
LinearRegression()
```

In []:

```
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```

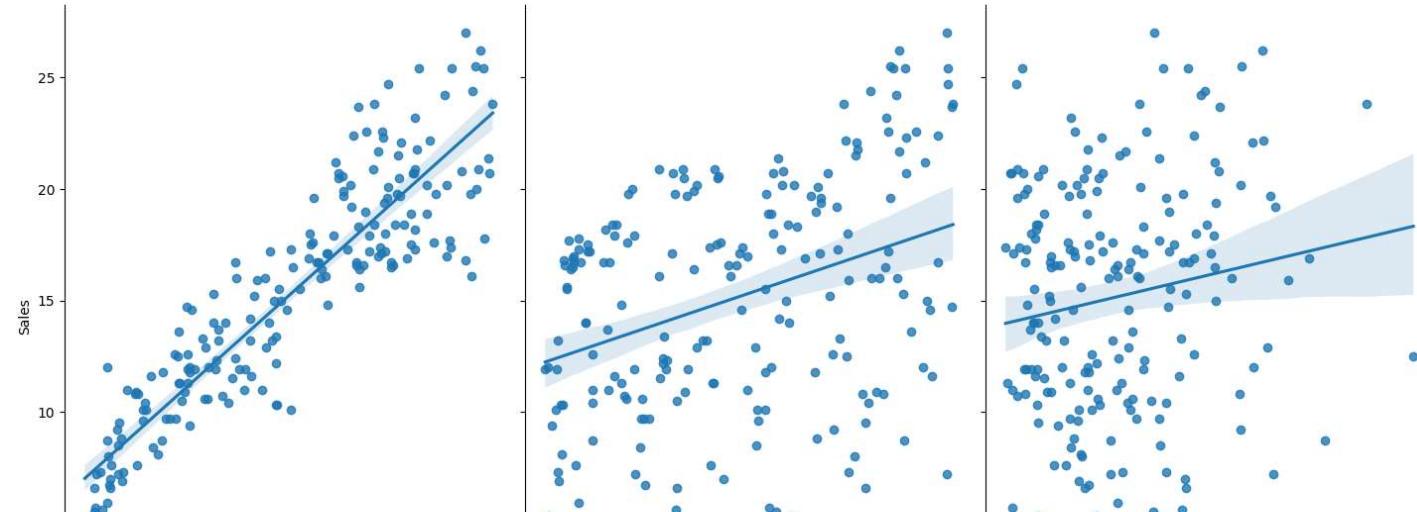


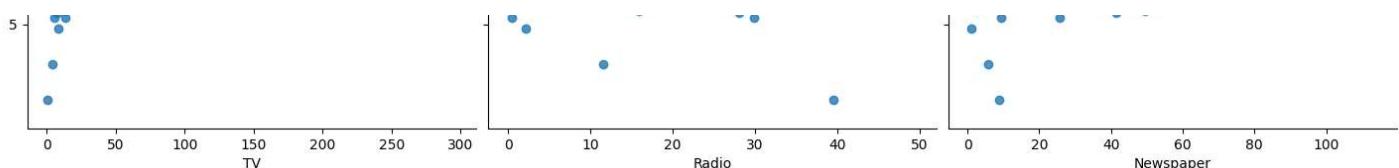
In []:

```
sns.pairplot(df,x_vars=['TV', 'Radio', 'Newspaper'],y_vars='Sales',height=7,aspect=0.7,k
ind='reg')
```

Out[]:

```
<seaborn.axisgrid.PairGrid at 0x7f0859087940>
```





In []:

```
#accuracy
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

0.6986549848574224

In []:

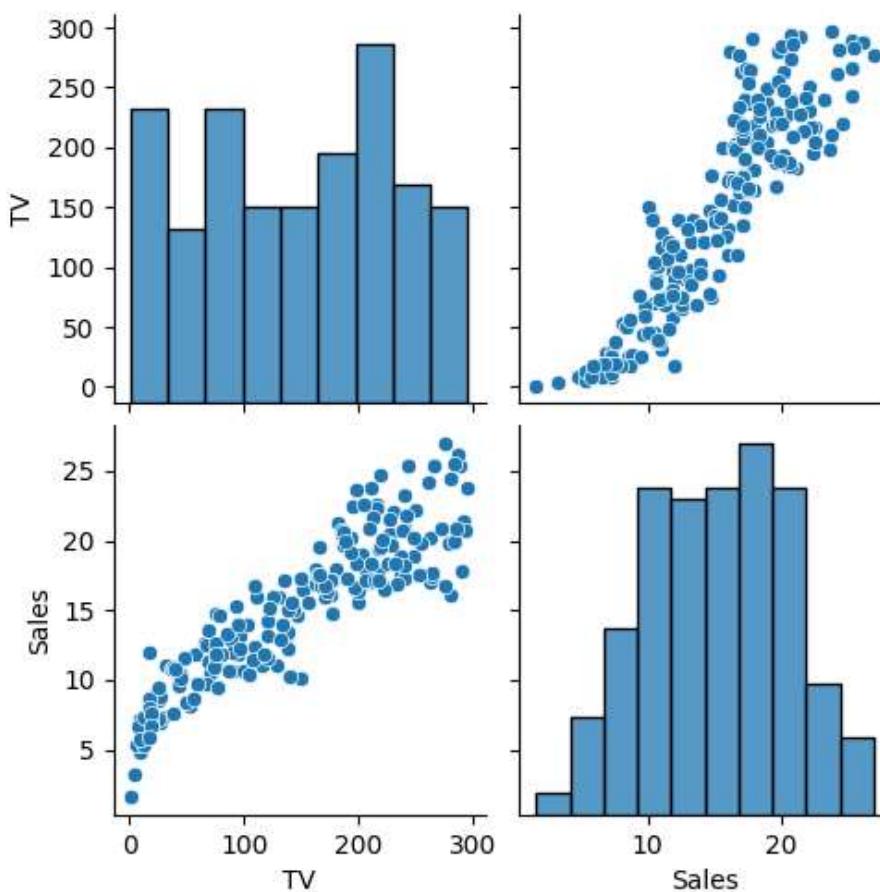
```
from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

In []:

```
ddf=df[['TV', 'Radio', 'Newspaper', 'Sales']]
```

In []:

```
df.drop(columns = ["Radio", "Newspaper"], inplace = True)
sns.pairplot(df)
df.Sales=np.log(df.Sales)
```



In []:

```
features=df.columns[0:2]
target=df.columns[-1]
X=df[features].values
y=df[target].values
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
```

```
scaler=StandardScaler()
X_train=scaler.fit_transform(X_train)
X_test=scaler.transform(X_test)
```

The dimension of X_train is (140, 2)
The dimension of X_test is (60, 2)

In []:

```
#Linear regression model
regr=LinearRegression()
regr.fit(X_train,y_train)
actual=y_test #actual value
train_score_regr=regr.score(X_train,y_train)
test_score_regr=regr.score(X_test,y_test)
print("\nLinear model:\n")
print("The train score for Linear model is {}".format(train_score_regr))
print("The test score for Linear model is {}".format(test_score_regr))
```

Linear model:

The train score for Linear model is 1.0
The test score for Linear model is 1.0

In []:

```
#ridge regression model
ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge model:

The train score for ridge model is 0.9902871391941609
The test score for ridge model is 0.9844266285141219

In []:

```
#using the linear cv model for ridge regression
from sklearn.linear_model import RidgeCV
#ridge cross validation
ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
#score
print(ridge_cv.score(X_train,y_train))
print(ridge_cv.score(X_test,y_test))
```

0.999999999997627
0.9999999999962467

In []:

```
#using the linear cv model for lasso regression
from sklearn.linear_model import LassoCV
#lasso cross validation
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y_train)
#score
print(lasso_cv.score(X_train,y_train))
print(lasso_cv.score(X_test,y_test))
```

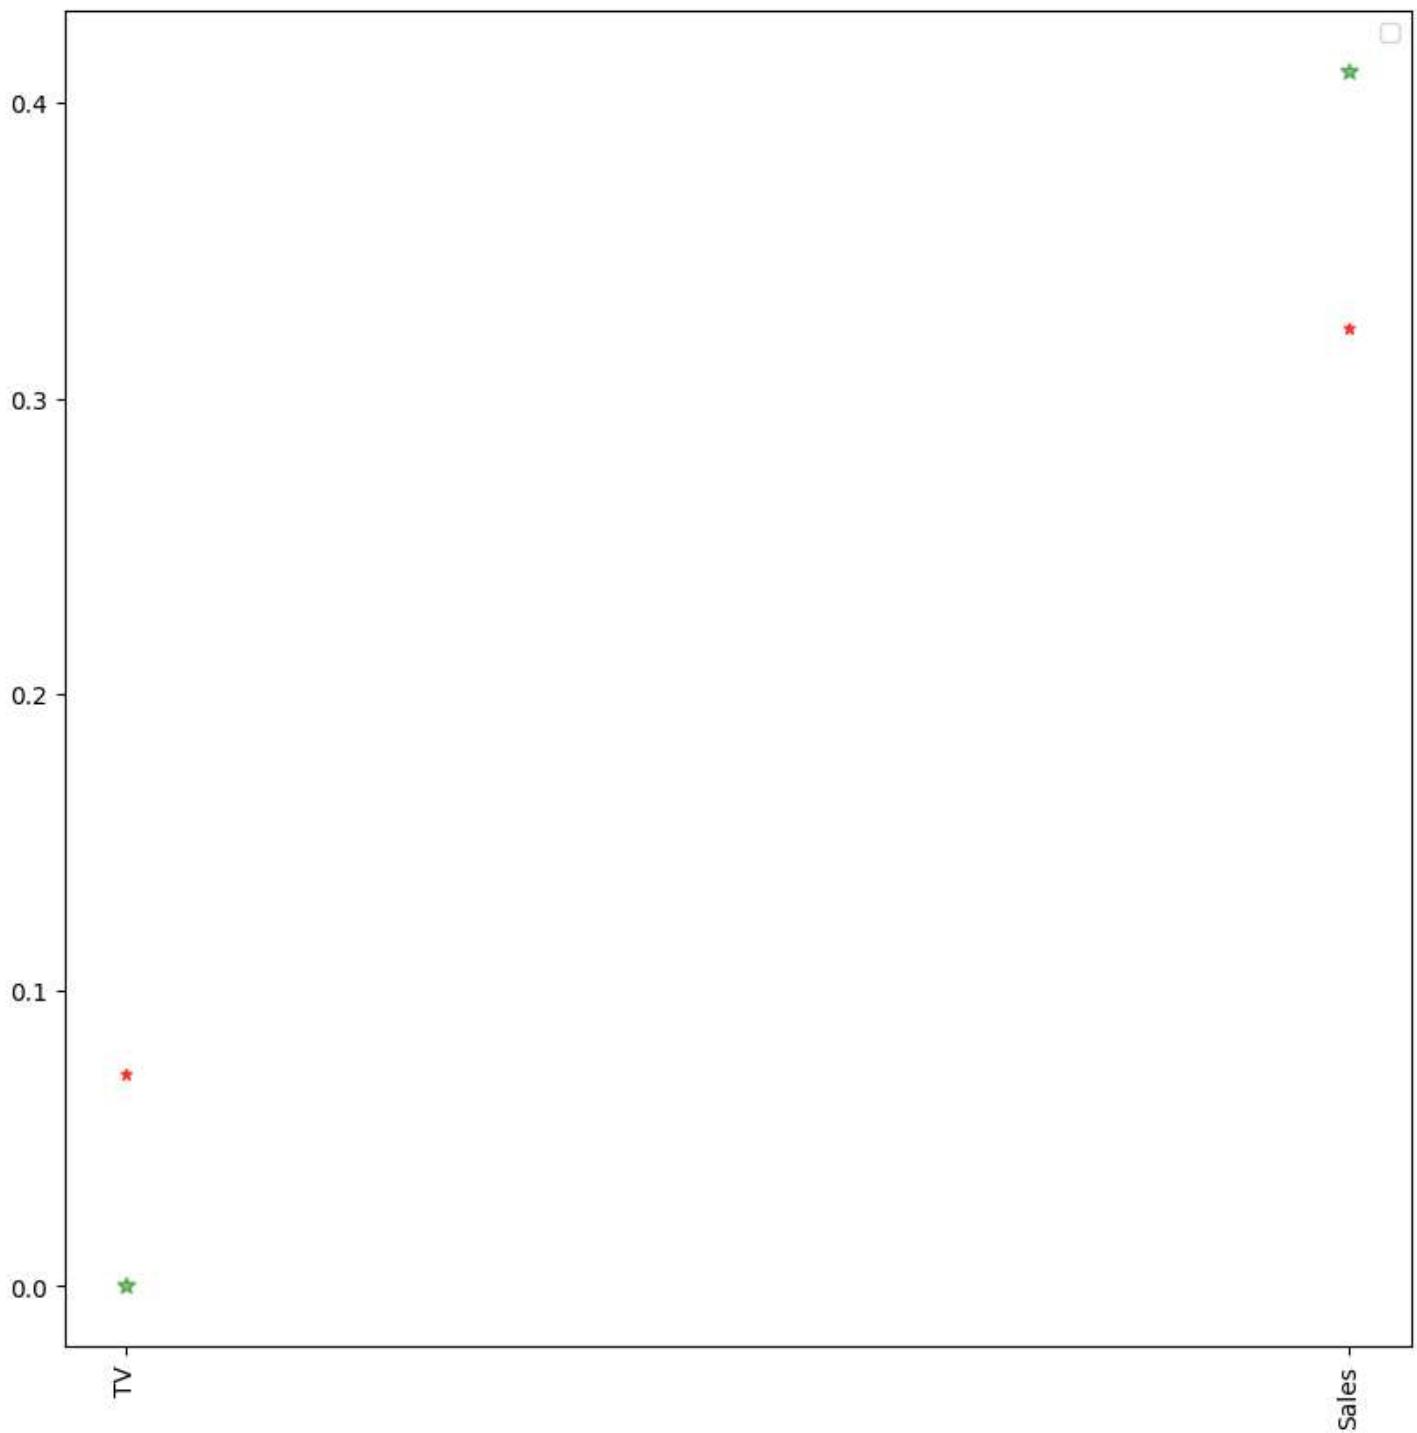
0.9999999343798134
0.9999999152638072

In []:

```
plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color
```

```
'red')
plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize=7,color='green')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



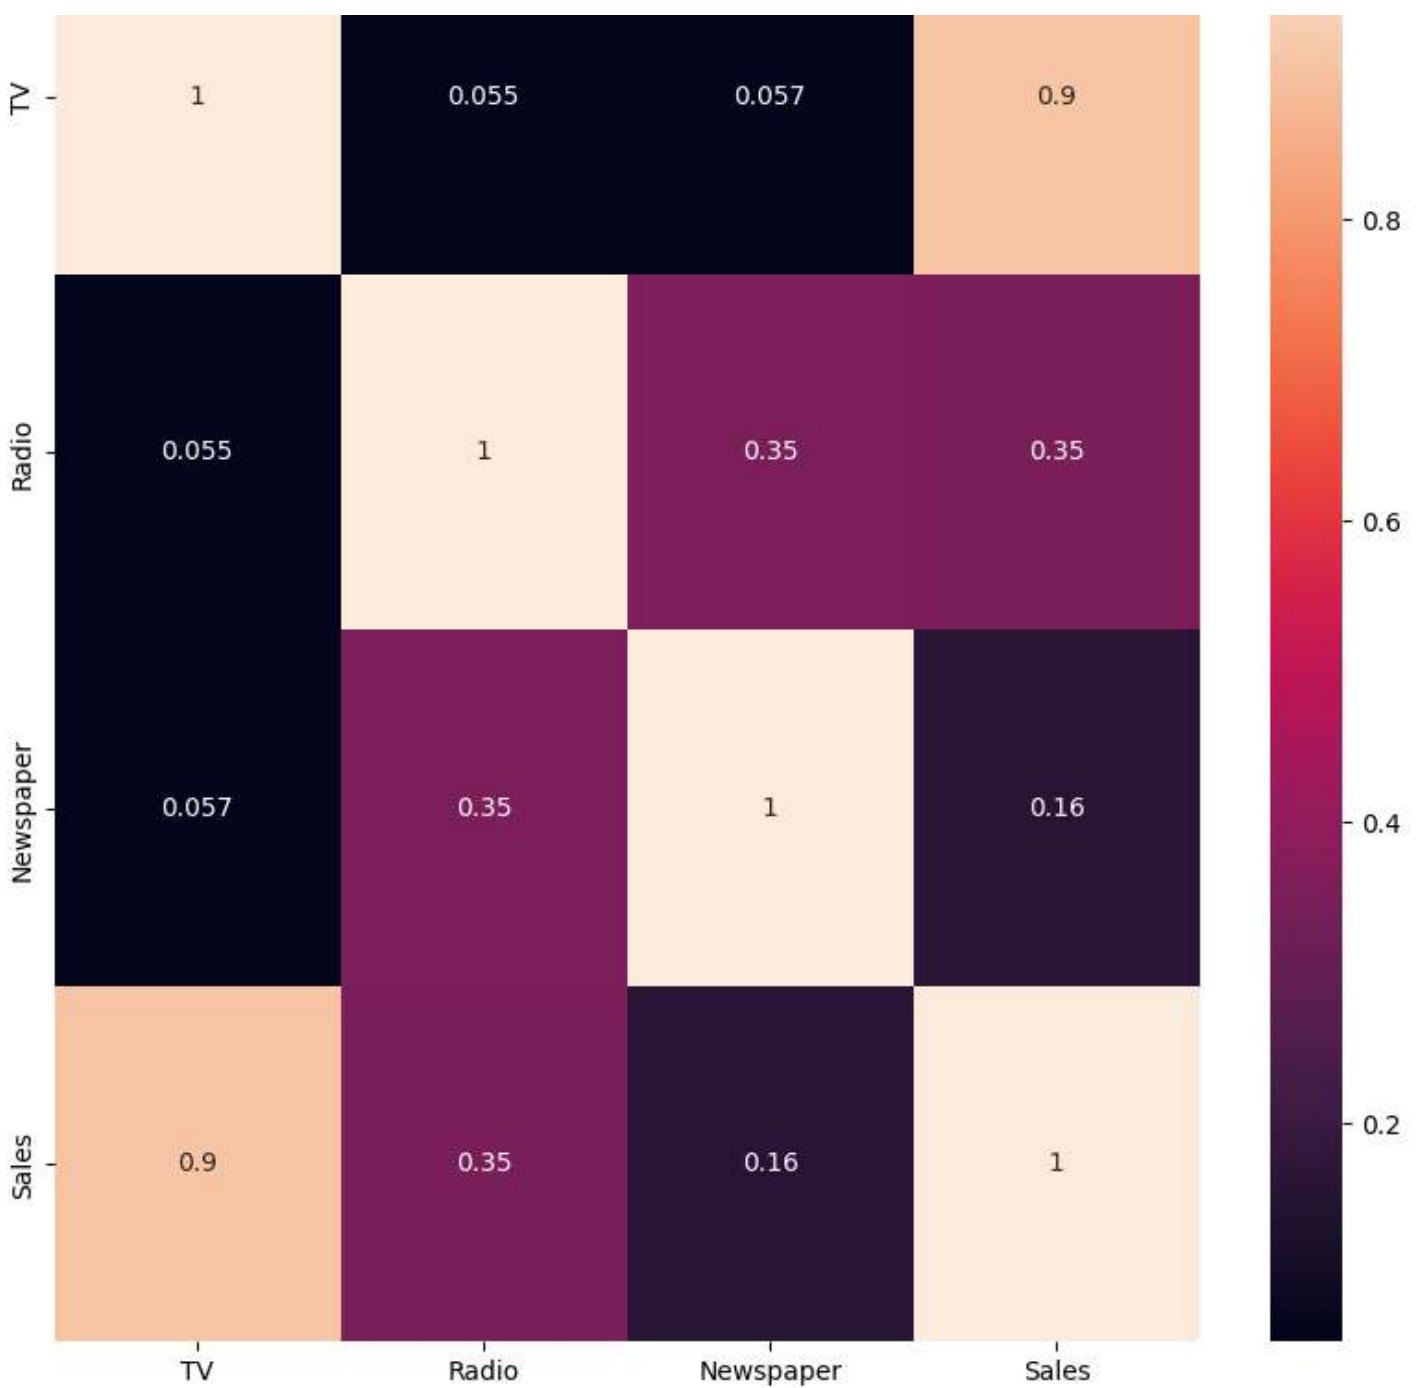
In []:

```
#ridge regression
plt.figure(figsize=(10,10))
sns.heatmap(ddf.corr(),annot=True)
```

Out[]:

<Axes: >





In []:

```
#lasso regression model
lassoReg=Lasso(alpha=10)
lassoReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_lasso=lassoReg.score(X_train,y_train)
test_score_lasso=lassoReg.score(X_test,y_test)
print("\nLasso model:\n")
print("The train score for lasso model is {}".format(train_score_lasso))
print("The test score for lasso model is {}".format(test_score_lasso))
```

Lasso model:

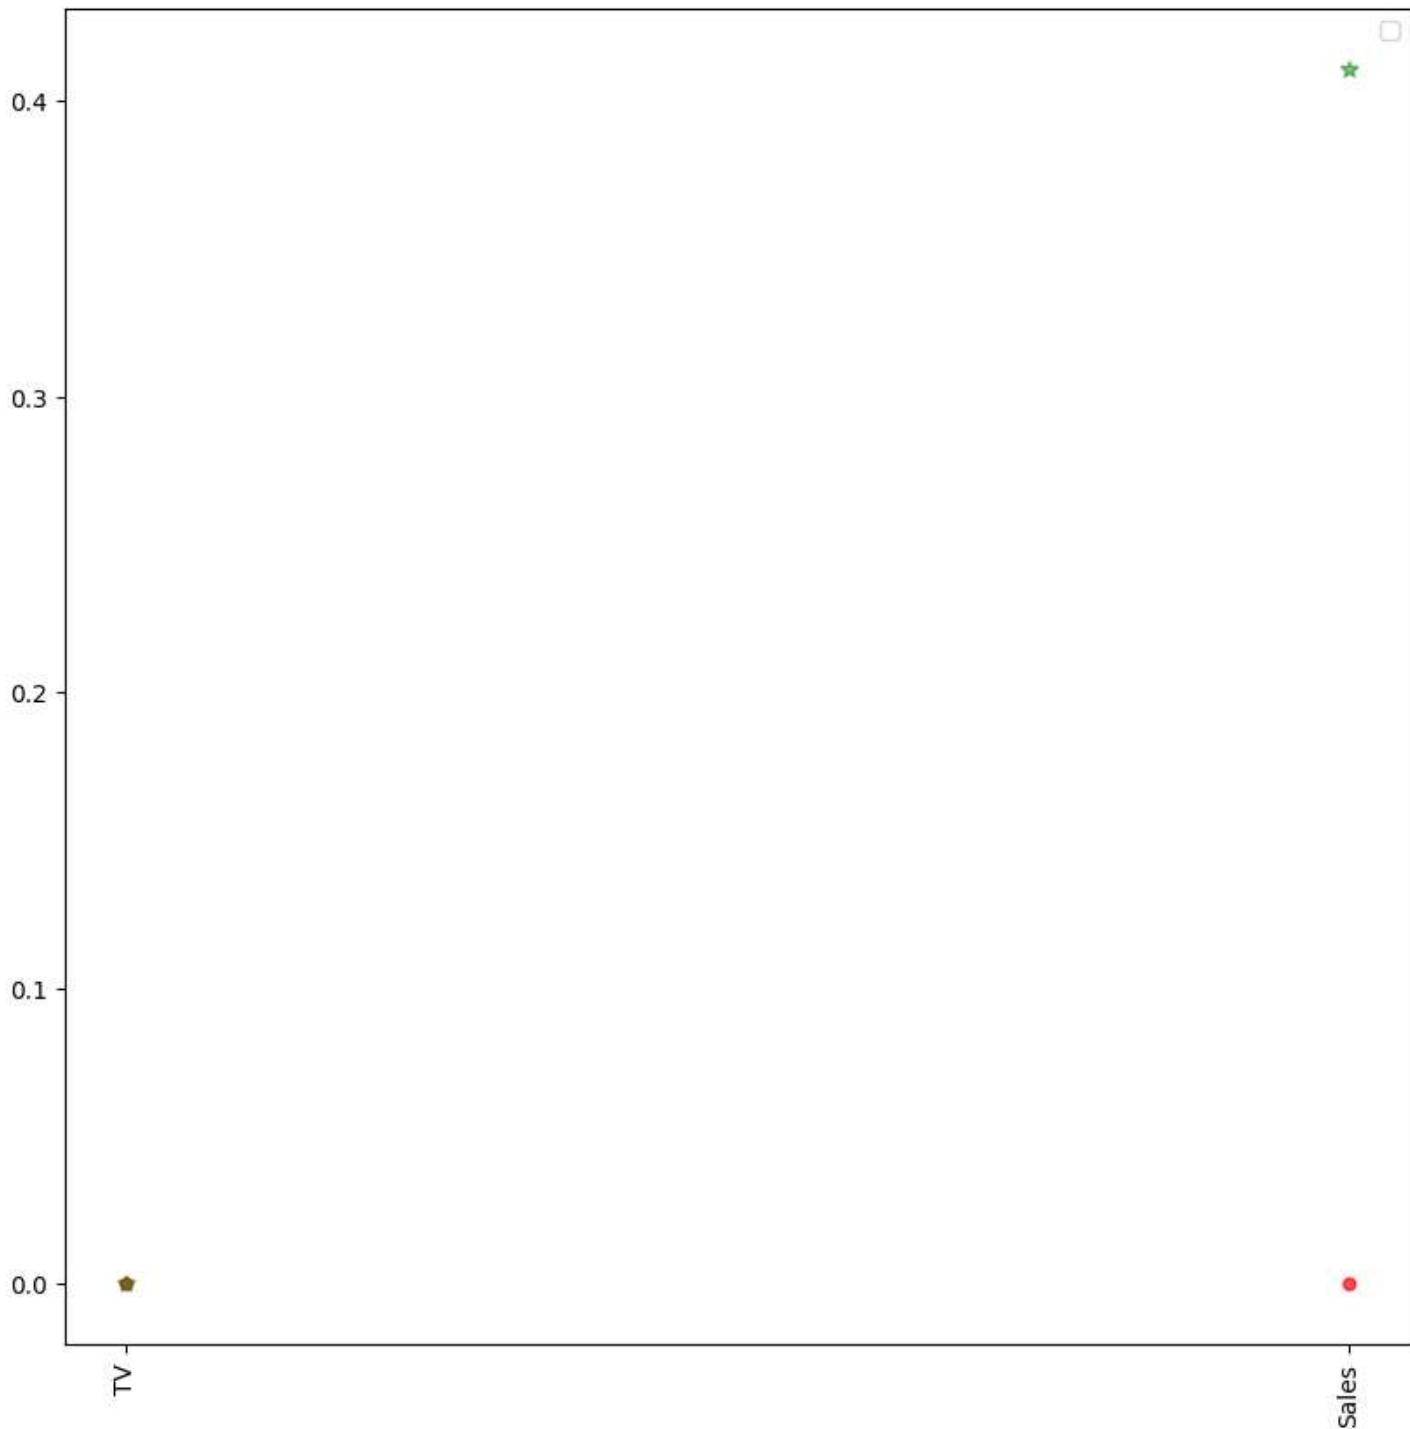
The train score for lasso model is 0.0
The test score for lasso model is -0.0042092253233847465

In []:

```
plt.figure(figsize=(10,10))
plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='o',markersize=5,color='red')
plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize=7,color='green')
plt.xticks(rotation=90)
```

```
plt.legend()  
plt.show()
```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



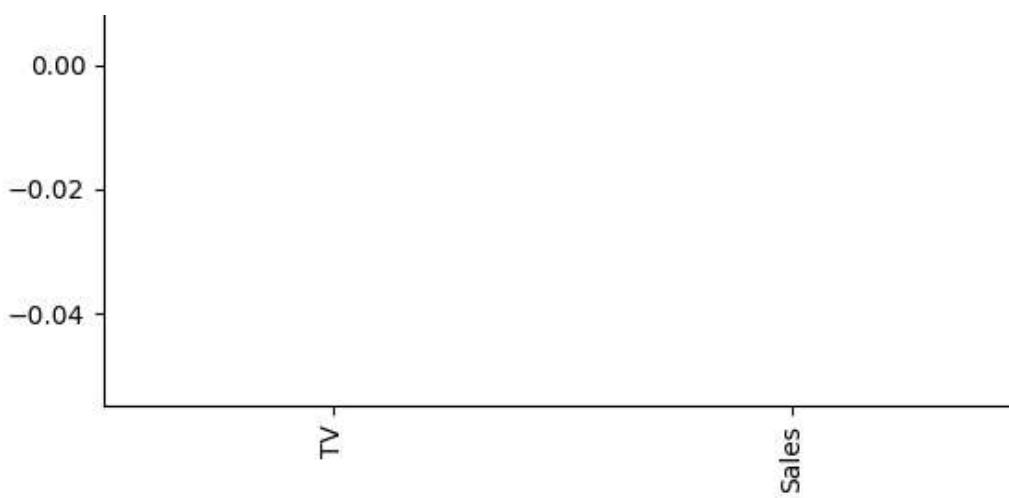
In []:

```
pd.Series(lassoReg.coef_, features).sort_values(ascending=True).plot(kind="bar")
```

Out[]:

```
<Axes: >
```



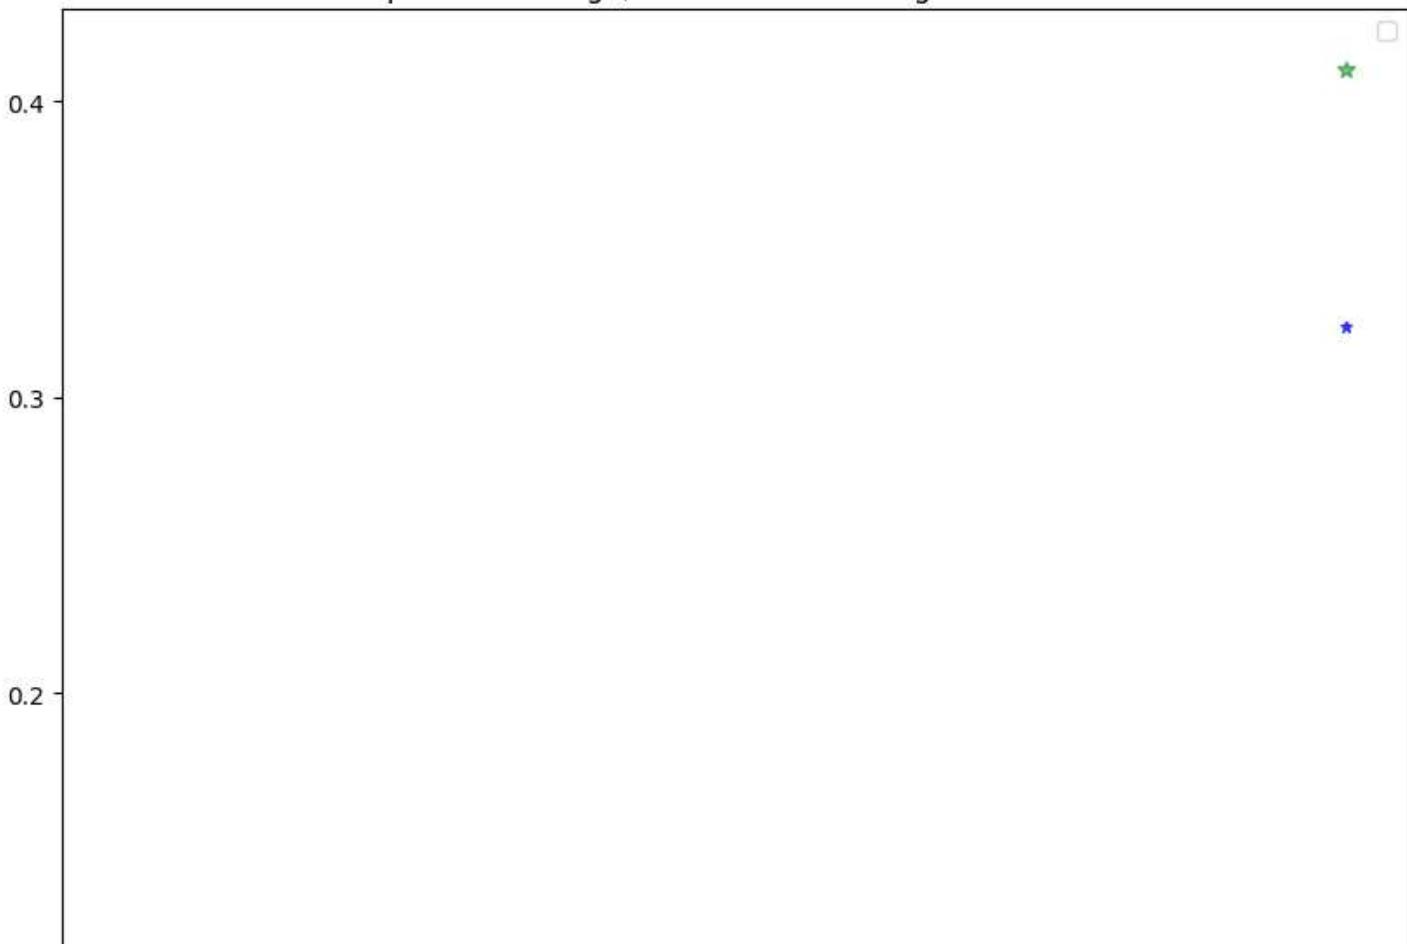


In []:

```
#plot size
plt.figure(figsize=(10,10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='blue')
#add plot for lasso regression
plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
#add plot for linear model
plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize=7,color='green')
#rotate axis
plt.xticks(rotation=90)
plt.legend()
plt.title("Comparison of Ridge,Lasso and Linear regression models")
plt.show()
```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

Comparison of Ridge,Lasso and Linear regression models





In []:

```
#elasticnet
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
y_pred_elastic=regr.predict(X_train)
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

[0.00417976 0.]
2.026383919311004
Mean Squared Error on test set 0.5538818050142158